

存储系统若干问题

刘绍辉

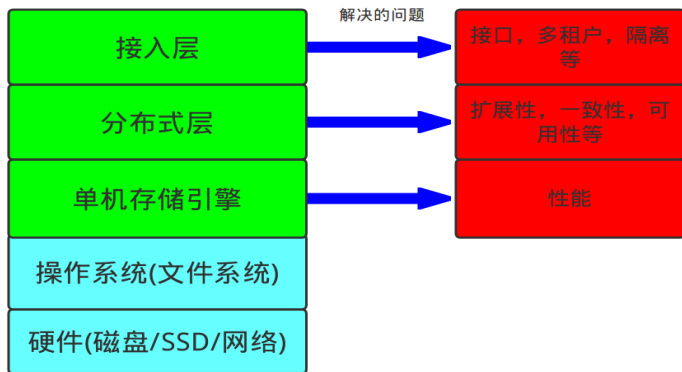
小米云存储组

May 20, 2016

目标

- ▶ 介绍存储系统基本层次和结构
- ▶ 核心问题取舍和实现方法

存储系统



提纲

1. 硬件
2. 操作系统
3. 单机存储引擎
4. 分布式层
5. 接入层

硬件

- ▶ 机械硬盘
- ▶ 固态硬盘(SSD)
- ▶ 内存(RAM is the new disk)
- ▶ 万兆网卡和网络拓扑



Traditional hard disk drive



Solid state hard drive

重要的数字

Numbers that every computer engineer should know

- ▶ Mutex lock/unlock: 17 ns
- ▶ Main memory reference: 100 ns
- ▶ Compress 1K bytes with Zippy: 10,000 ns
- ▶ SSD random read: 16 μ s
- ▶ Hard Disk seek: 8 ms
- ▶ Round trip within same datacenter: 300 μ s
- ▶ Round trip across zones in same region: 800 μ s
- ▶ Read 1 MB sequentially from memory: 12 μ s
- ▶ Read 1 MB sequentially from SSD: 200 μ s
- ▶ Read 1 MB sequentially from disk: 2ms

http://www.eecs.berkeley.edu/~rcs/research/interactive_latency.html

测试工具

- ▶ 磁盘：fio
- ▶ 网络：ping/iperf
- ▶ mysql工具：pt-mysql-summary

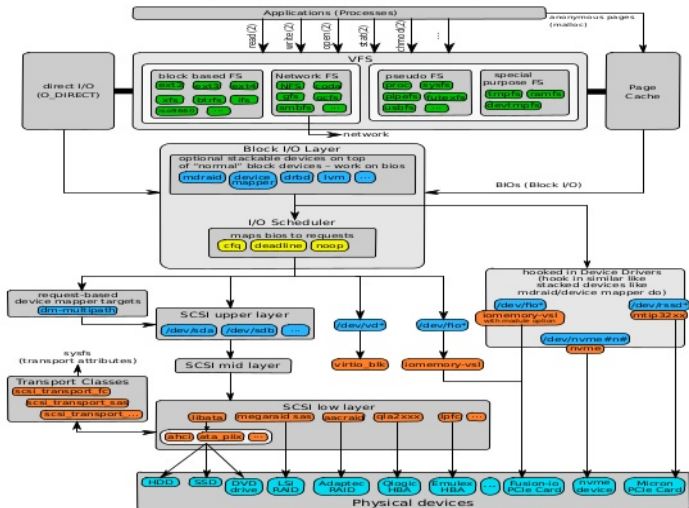
提纲

1. 硬件
2. 操作系统
3. 单机存储引擎
4. 分布式层
5. 接入层

文件系统

The Linux I/O Stack Diagram

version 0.1, 2012-03-06
outlines the Linux I/O stack as of Kernel version 3.3



The Linux I/O Stack Diagram (version 0.1, 2012-03-06)
http://www.kernel.org/doc/Documentation/iostack-diagram.html
Copyright © 2012-2013, Red Hat, Inc. All rights reserved.
License: GPL v2 or later, http://www.gnu.org/licenses/gpl-2.0.html

数据持久性

- ▶ write: Only write data to OS page cache
- ▶ sync: flush all dirty buffers to disk
- ▶ fsync: flush all blocks that belong to a specific open file to disk
- ▶ fdatasync: flush all data blocks(no inode block) that belong to a specific open file to disk

观测工具

- ▶ 磁盘：iostat/iotop
- ▶ 网络：iftop
- ▶ sysdig systemtap

吐血推荐：Red Hat Enterprise Linux Performance Tuning Guide

提纲

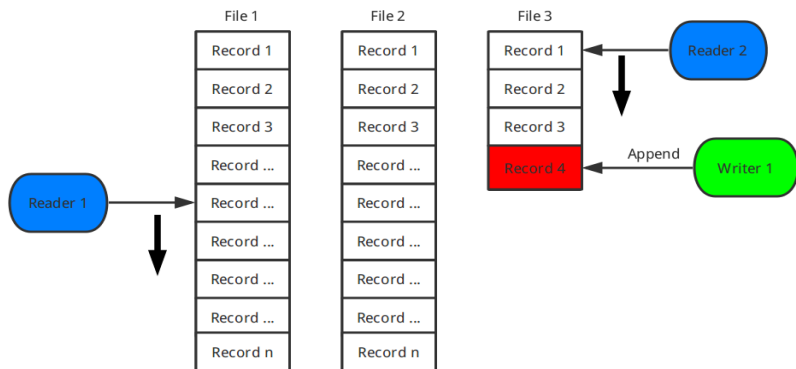
1. 硬件
2. 操作系统
3. 单机存储引擎
4. 分布式层
5. 接入层

单机存储引擎

- ▶ 文件: HDFS/Kafka datanode
- ▶ Hash Table (BitCask/Memcached/Redis)
- ▶ B+ Tree (MySQL innoDB engine)
- ▶ LSM Tree (LevelDB/RocksDB)

<http://www.xaprb.com/blog/2015/04/02/state-of-the-storage-engine/>

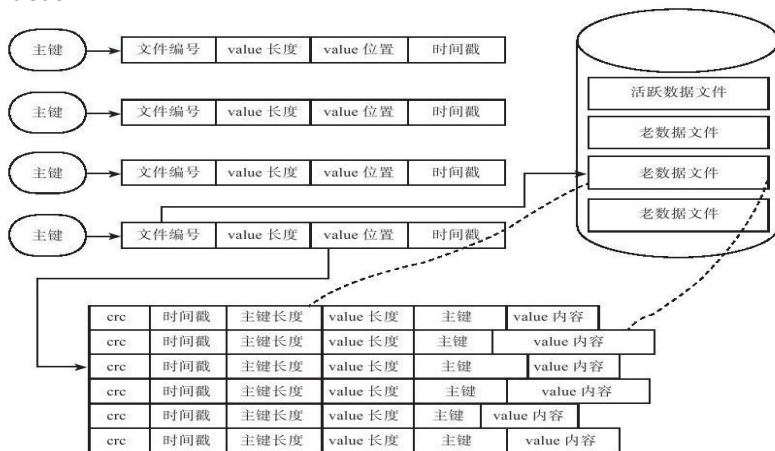
File: 顺序读写



HDFS/Kafka datanode

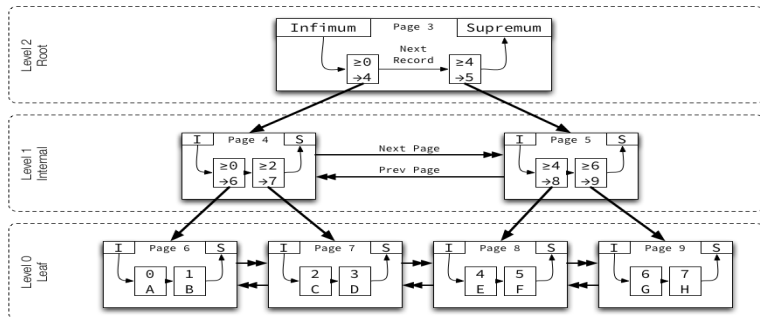
Hash Table: 随机读写

BitCask



面向小文件

B+Tree Structure



Levels are numbered starting from 0 at the leaf pages, incrementing up the tree.

Pages on each level are doubly-linked with previous and next pointers in ascending order by key.

Records within a page are singly-linked with a next pointer in ascending order by key.

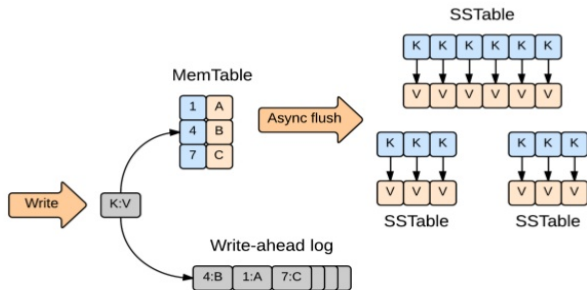
Infimum represents a value lower than any key on the page, and is always the first record in the singly-linked list of records.

Supremum represents a value higher than any key on the page, and is always the last record in the singly-linked list of records.

Non-leaf pages contain the minimum key of the child page and the child page number, called a "node pointer".

面向读

Log-structured merge tree layout



面向写，问题：写放大，标记删除

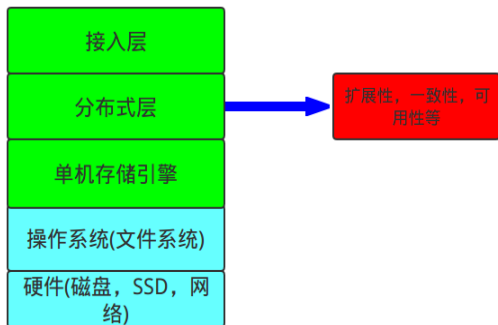
单机存储引擎

为什么需要这么多不同存储引擎？
是否有一个引擎满足所有需求？

提纲

1. 硬件
2. 操作系统
3. 单机存储引擎
4. 分布式层
5. 接入层


分布式层



扩展性

数据

- ▶ File : (Block1, Block2, ...)
- ▶ Table : (Range1, Range2, ...) / (Hash1, Hash2, ...)
- ▶ Queue : (partition1, partition2, ...)





扩展性

元数据

- ▶ partition信息
- ▶ 路由信息等等

中心化(Master-Slave) vs 去中心化(P2P)



Cassandra concepts, patterns and anti-patterns - ApacheCon EU 2012

Master-Worker优势：结构简单，容易控制
问题：单点问题

一致性和可用性


在保证强一致的情况下如何保证更高的可用性。

一致性分类：


- ▶ 强一致：写成功返回之后马上可以读到最新结果
- ▶ 弱一致：最终一致性/因果一致性/顺序一致性等等

一致性和可用性

- ▶ Primary-backup协议(Zab协议)
- ▶ state machine replication(Paxos/Raft/Pacific)



Primary Backup



State Machine Replication

Paxos

核心：尊重之前的决定，并帮忙他完成

```
## Paxos Read and Write

### Phase 1:
prepare(paxos_intance, proposal_id) {
    if proposal_id <= paxos_intance.promised_proposal_id:
        return
    paxos_intance.promised_proposal_id = proposal_id
    if paxos_intance.accepted_proposal_id == -1:
        return promise();
    else
        return promise(paxos_intance.accepted_proposal_id, paxos_intance.accepted_proposal_value);
}

### Phase 2:
proposal_value = accepted_proposal_value with highest accepted_proposal_id or any value.

accept(paxos_intance, proposal_id, proposal_value)
    if proposal_id >= paxos_intance.promised_proposal_id:
        paxos_intance.accepted_proposal_id = proposal_id;
        paxos_intance.accepted_proposal_value = proposal_value;
        return accepted()
```

其他问题

- ▶ Master的高可用(Paxos协议/Zookeeper + HDFS)
- ▶ 跨机房备份和容灾
- ▶ 多机房服务
- ▶ 分布式事务

提纲

1. 硬件
2. 操作系统
3. 单机存储引擎
4. 分布式层
5. 接入层


接入层

服务化

- ▶ 接口（给用户的承诺）：
File/Queue/Key-Value/Table/SQL/Key-Objects
- ▶ rpc or rest
- ▶ 多租户：用户认证和权限检查
- ▶ 隔离：Quota限制
- ▶ 审计和计费

理想状态

- ▶ Master-Worker架构
- ▶ 水平扩展，Shard策略可配置
- ▶ 类Paxos副本复制协议
- ▶ 支持多种存储引擎
- ▶ 接口：File/Queue/Key-Value/Table/SQL/Key-Objects
- ▶ 多租户和隔离



参考

- ▶ MapReduce: Simplified Data Processing on Large Clusters
- ▶ The Google File System
- ▶ Bigtable: A Distributed Storage System for Structured Data
- ▶ Megastore: Providing Scalable, Highly Available Storage for Interactive Services
- ▶ Spanner: Google's Globally-Distributed Database
- ▶ F1 - The Fault-Tolerant Distributed RDBMS Supporting Google's Ad Business

开源项目

- ▶ TiDB : <https://github.com/pingcap/tidb>
- ▶ cockroachdb: <https://github.com/cockroachdb/cockroach>

图书

- ▶ <<大规模分布式存储系统：原理解析与架构实战>> 杨传辉
- ▶ <<Distributed systems: for fun and profit>>
<https://github.com/mixu/distsysbook>